

Explain the distinctions among the terms primary key, candidate key, and superkey.

Keys are used to uniquely identify rows within data. A superkey is any set of columns that, when put together, uniquely identifies every row in the table. There are no two distinct rows, or tuples, that have the same values for the attributes in the set. When looking for the superkey to identify every row in the table, you must consider all columns in the table. Stemming from the superkey, a candidate key is an identifier that uses the smallest number of columns to achieve uniqueness of every row in the column. A candidate key cannot have any columns removed from it without losing the unique identification property. From a candidate key, we get a primary key. A primary key is the chosen candidate key to achieve uniqueness throughout the rows when you have multiple candidate keys. Out of the three keys listed, a primary key is the smallest identifier of uniqueness.

Write a short essay on data types. Select a topic for which you might create a table. Name the table and list its fields (columns). For each field, give its data type and whether or not it is nullable.

There are many kinds of data types within a relation schema in SQL. Data types include character strings, bit strings, Boolean data, integers, floating-point numbers, and dates and times. Character strings can be of either fixed or varying length. Fixed-length character strings are denoted by the term CHAR, whereas strings of “x” amount of characters are denoted by VARCHAR in SQL. Bit strings are similar to characters, except their values are in bits. BIT denotes the fixed amount of bits, whereas BITVARYING means that there is a varying length of bits. Boolean data is a type whose value is logical. For example, Boolean data would be TRUE, FALSE, or UNKNOWN. Integer (INT or INTEGER) is any typical integer value. Within INTEGER, there is SHORTINT, which also denotes an integer, but there are less numbers of bits permitted. As for floating numbers, they are denoted by the term FLOAT, REAL or DECIMAL. For more exact floating numbers, there is DOUBLE PRECISION and DECIMAL (n,d). DECIMAL (n,d) allows you to exactly determine what the number will look like; n is the amount of real numbers, and d represents the positions from the right of the decimal point. Finally, dates and times are denoted by DATE and TIME. Because dates and times require a special form, these data types are strung together in a unique way to create that special form that makes sense as a date or time.

A topic that serves as a good example of data types within a table would be a table of Employee Profiles for an organization. This table would keep all the information about the company’s employees, including everything from their personal information to their organizational information. Some of the fields of the table, for example, would include employee

first name (VARCHAR), employee last name (VARCHAR), middle initial (CHAR), age (INT), department (VARCHAR), position (CHAR or VARCHAR), date of birth (DATE), marital status (Boolean TRUE, FALSE or UNKNOWN), salary (fixed salary with INT, hourly wage with DECIMAL or DOUBLE PRECISION). For all of these fields, I do not think that there are any that are nullable. There is an answer for each field, even if the answer is unknown (for example, marital status). I believe the only time these fields can be nullable is when an employee is initially entered into the system (for example, a new hire). It may take time to get the paperwork that has the information on the employee's profile, so in the short time that there is no information, "nullable" is acceptable.

Explain the following relational "rules" with examples and reasons why they are important.

The "first normal form" rule

The first relational rule states that there should be no multi-valued attributes in the fields of a table. This means that there cannot be more than one data type in a column. If there needs to be a variety of data types associated with a field, then there needs to be a cross reference between different tables in the database. Information can only be represented in one way in a table. Going off of my previous example of Employee Profiles, one of the fields could be the number of dependents. In this field, there would be a number that would then reference another table listing the specifics of each dependent. The field of "number of dependents" in the Employee Profile table cannot list a various amount of data types, such as each name of the dependent. If it did that, then it would not be able to list all the information for each dependent.

The "access rows by content only" rule

The second relation rule states that, when asking about data, ask "what, not where." This rule is especially important because by specifying *what* data you're asking for, rather than the location of the data (*where* it is), you are requesting specifically what you are searching for. If you ask *where* (for example, "for the information in column three"), then there is a chance that the information in column three has changed.

The "all rows must be unique" rule

The third relation rule states that all rows in a table must be unique. For each row in a table, every field is in the row. If there is repeated information in the rows, then you cannot specify *what* information you need (this relates back to the previous rule). For example, if there are two rows within the Employee Profile table with the same information (same first name, same last name, same department, etc.), then there is no way to tell them apart. This is the reason for keys – we need to uniquely identify rows from one another in tables.

pgAdmin3 File Edit Query Favourites Macros View Window Help

Query - postgres on postgres@localhost:5432 *

postgres on postgres@localhost:5432

SQL Editor Graphical Query Builder

Previous queries

```
select *  
from products;  
  
select *  
from orders;
```

Output pane

Data Output Explain Messages History

	ordno integer	mon character(3)	cid character(4)	aid character(3)	pid character(3)	qty integer	dollars numeric(12,2)
1	1011	jan	c001	a01	p01	1000	450.00
2	1013	jan	c002	a03	p03	1000	880.00
3	1015	jan	c003	a03	p05	1200	1104.00
4	1016	jan	c006	a01	p01	1000	500.00
5	1017	feb	c001	a06	p03	600	540.00
6	1018	feb	c001	a03	p04	600	540.00
7	1019	feb	c001	a02	p02	400	180.00
8	1020	feb	c006	a03	p07	600	600.00
9	1021	feb	c004	a06	p01	1000	460.00
10	1022	mar	c001	a05	p06	400	720.00
11	1023	mar	c001	a04	p05	500	450.00
12	1024	mar	c006	a06	p01	800	400.00
13	1025	apr	c001	a05	p07	800	720.00
14	1026	may	c002	a05	p03	800	740.00

OK. Unix Ln 197, Col 1, Ch 5884 21 chars 14 rows. 19 ms

pgAdmin3 File Edit Query Favourites Macros View Window Help

Query - postgres on postgres@localhost:5432 *

postgres on postgres@localhost:5432

SQL Editor Graphical Query Builder

Previous queries Delete Delete All

```
-- Connect to your Postgres server and set the active database to CAP2. Then . . .  
  
select *  
from customers;  
  
select *  
from agents;  
  
select *  
from products;  
  
select *  
from orders;
```

Output pane

Data Output Explain Messages History

	pid character(3)	name text	city text	quantity integer	priceusd numeric(10,2)
1	p01	comb	Dallas	111400	0.50
2	p02	brush	Newark	203000	0.50
3	p03	razor	Duluth	150600	1.00
4	p04	pen	Duluth	125300	1.00
5	p05	pencil	Dallas	221400	1.00
6	p06	folder	Dallas	123100	2.00
7	p07	case	Newark	100500	1.00
8	p08	clip	Newark	200600	1.25

OK. Unix Ln 194, Col 1, Ch 5859 23 chars 8 rows. 13 ms

pgAdmin3 File Edit Query Favourites Macros View Window Help

Query - postgres on postgres@localhost:5432 *

postgres on postgres@localhost:5432

SQL Editor Graphical Query Builder

Previous queries Delete Delete All

```
-- Connect to your Postgres server and set the active database to CAP2. Then . . .  
  
select *  
from customers;  
  
select *  
from agents;  
  
select *  
from products;  
  
select *  
from orders;
```

Output pane

Data Output Explain Messages History

	aid character(3)	name text	city text	percent real
1	a01	Smith	New York	6
2	a02	Jones	Newark	6
3	a03	Brown	Tokyo	7
4	a04	Gray	New York	6
5	a05	Otasi	Duluth	5
6	a06	Smith	Dallas	5
7	a08	Bond	London	7

OK. Unix Ln 191, Col 1, Ch 5836 21 chars 7 rows. 18 ms

pgAdmin3 File Edit Query Favourites Macros View Window Help

Query - postgres on postgres@localhost:5432 *

postgres on postgres@localhost:5432

SQL Editor Graphical Query Builder

Previous queries Delete Delete All

```
VALUES(1025, 'apr', 'c001', 'a05', 'p07', 800, 720.00);

INSERT INTO Orders( ordno, mon, cid, aid, pid, qty, dollars )
VALUES(1026, 'may', 'c002', 'a05', 'p03', 800, 740.00);

-- SQL statements for displaying example data into the CAP2 database
-- Connect to your Postgres server and set the active database to CAP2. Then . . .

select *
from customers;

select *
from agents;
```

Output pane

Data Output Explain Messages History

	cid character(4)	name text	city text	discount numeric(5,2)
1	c001	Tiptop	Duluth	10.00
2	c002	Basics	Dallas	12.00
3	c003	Allied	Dallas	8.00
4	c004	ACME	Duluth	8.00
5	c005	Weyland-Yutani	Acheron	0.00
6	c006	ACME	Kyoto	0.00

OK. Unix Ln 188, Col 1, Ch 5810 24 chars 6 rows. 16 ms